

Off-policy methods with approximations

Master Degree in Control Systems Engineering 2023-2024
University of Padova

Teachers: Prof. Ruggero Carli & Prof. Gian Antonio Susto

Today's lecture

- Why off-policy methods do not work (counterexamples + deadly triad)
- Linear value-function geometry – TD gradient (only mentioned)

Off-policy methods

π : **target policy**

b : **exploring policy**

in off-policy learning we seek to learn a value function for a target policy π , given data due to a different behavior policy b

Off-policy methods

π : **target policy**

b : **exploring policy**

in off-policy learning we seek to learn a value function for a target policy π , given data due to a different behavior policy b

Many of the tabular off-policy algorithms use the *importance sampling ratio* :

$$\rho_t \doteq \rho_{t:t} = \frac{\pi(A_t|S_t)}{b(A_t|S_t)}$$

$$\rho_{t:T-1} \doteq \frac{\prod_{k=t}^{T-1} \pi(A_k|S_k)p(S_{k+1}|S_k, A_k)}{\prod_{k=t}^{T-1} b(A_k|S_k)p(S_{k+1}|S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$$

$$\rho_{t:h} \doteq \prod_{k=t}^{\min(h, T-1)} \frac{\pi(A_k|S_k)}{b(A_k|S_k)}$$

Off-policy methods with function Approx. : semi-gradient alg.

Natural extension : *To convert them to semi-gradient form, we simply replace the update to an array (V or Q) to an update to a weight vector ($\mathbf{w}$), using the approximate value function ($\hat{v}$ or $\hat{q}$) and its gradient.*

Off-policy methods with function Approx. : semi-gradient alg.

Natural extension : To convert them to semi-gradient form, we simply replace the update to an array (V or Q) to an update to a weight vector ($\mathbf{w}$), using the approximate value function ($\hat{v}$ or $\hat{q}$) and its gradient.

Examples : the one-step, state-value, semi-gradient off-policy TD(0)

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, \mathbf{w}_t)$$

where

$$\delta_t \doteq R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t)$$

and

$$\rho_t \doteq \rho_{t:t} = \frac{\pi(A_t | S_t)}{b(A_t | S_t)}$$

Off-policy methods with function Approx. : semi-gradient alg.

Natural extension : To convert them to semi-gradient form, we simply replace the update to an array (V or Q) to an update to a weight vector ($\mathbf{w}$), using the approximate value function ($\hat{v}$ or $\hat{q}$) and its gradient.

Examples : the one-step, action values, semi-gradient Expected Sarsa

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \delta_t \nabla \hat{q}(S_t, A_t, \mathbf{w}_t)$$

where

$$\delta_t \doteq R_{t+1} + \gamma \sum_a \pi(a|S_{t+1}) \hat{q}(S_{t+1}, a, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t)$$

Notice: no use of importance sampling

Off-policy methods with function Approx. : semi-gradient alg.

These algorithms do not converge as robustly as they do under on-policy training

Off-policy methods with function Approx. : semi-gradient alg.

These algorithms do not converge as robustly as they do under on-policy training

The extension to function approximation turns out to be significantly different and harder for off-policy learning than it is for on-policy learning

Off-policy methods with function Approx. : semi-gradient alg.

These algorithms do not converge as robustly as they do under on-policy training

The extension to function approximation turns out to be significantly different and harder for off-policy learning than it is for on-policy learning

An issue : distribution of the updates

Off-policy methods with function Approx. : semi-gradient alg.

These algorithms do not converge as robustly as they do under on-policy training

The extension to function approximation turns out to be significantly different and harder for off-policy learning than it is for on-policy learning

An issue : distribution of the updates



Distribution of updates in the off-policy case is not according to the on-policy distribution

Off-policy methods with function Approx. : semi-gradient alg.

These algorithms do not converge as robustly as they do under on-policy training

The extension to function approximation turns out to be significantly different and harder for off-policy learning than it is for on-policy learning

An issue : distribution of the updates



Distribution of updates in the off-policy case is not according to the on-policy distribution

Remember : The on-policy distribution is important to the stability of semi-gradient methods.


Examples of Off-policy Divergence

To establish intuitions, we start by considering a very simple example.

Examples of Off-policy Divergence

To establish intuitions, we start by considering a very simple example.

Suppose, as part of a larger MDP, there are two states whose estimated values are of the functional form w and $2w$, where the parameter vector w consists of only a single component w .




expressions inside the two circles indicate the two state's values

Examples of Off-policy Divergence

To establish intuitions, we start by considering a very simple example.

Suppose, as part of a larger MDP, there are two states whose estimated values are of the functional form w and $2w$, where the parameter vector w consists of only a single component w .




expressions inside the
two circles indicate the
two state's values

This occurs under linear function approximation if the feature vectors for the two states are each simple numbers (single-component vectors), in this case 1 and 2.

Examples of Off-policy Divergence

To establish intuitions, we start by considering a very simple example.

Suppose, as part of a larger MDP, there are two states whose estimated values are of the functional form w and $2w$, where the parameter vector w consists of only a single component w .



This occurs under linear function approximation if the feature vectors for the two states are each simple numbers (single-component vectors), in this case 1 and 2.


Only one action is possible in this state : a deterministic transition to state 2 with reward 0

Examples of Off-policy Divergence

Suppose initially $w = 10$



The transition will then be from a state of estimated value 10 to a state of estimated value 20.




Examples of Off-policy Divergence

Suppose initially $w = 10$



The transition will then be from a state of estimated value 10 to a state of estimated value 20.



$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t) :$$


$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, \mathbf{w}_t)$$

If γ is nearly 1, then the TD error will be nearly 10, and, if $\alpha = 0.1$, then w will be increased to nearly 11 in trying to reduce the TD error.

Examples of Off-policy Divergence

Suppose initially $w = 10$

↳ The transition will then be from a state of estimated value 10 to a state of estimated value 20.




However, the second state's estimated value will also be increased, to nearly 22

Examples of Off-policy Divergence

Suppose initially $w = 10$



The transition will then be from a state of estimated value 10 to a state of estimated value 20.



However, the second state's estimated value will also be increased, to nearly 22


If the transition occurs again, then it will be from a state of estimated value ≈ 11 to a state of estimated value ≈ 22 , for a TD error of ≈ 11 - **larger, not smaller, than before.**

Examples of Off-policy Divergence

Suppose initially $w = 10$



The transition will then be from a state of estimated value 10 to a state of estimated value 20.




However, the second state's estimated value will also be increased, to nearly 22

If the transition occurs again, then it will be from a state of estimated value ≈ 11 to a state of estimated value ≈ 22 , for a TD error of ≈ 11 - **larger, not smaller, than before.**

It will look even more like the first state is undervalued, and its value will be increased again, this time to ≈ 12.1 . This looks bad, and in fact with further updates w will diverge to infinity.

Examples of Off-policy Divergence

Let us see more carefully what seen in the previous slide at the sequence of updates




The TD error on a transition between the two states is

$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t) = 0 + \gamma 2w_t - w_t = (2\gamma - 1)w_t$$

Examples of Off-policy Divergence

Let us see more carefully what seen in the previous slide at the sequence of updates



The TD error on a transition between the two states is


$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t) = 0 + \gamma 2w_t - w_t = (2\gamma - 1)w_t$$

and the off-policy semi-gradient TD(0) update is

$$\begin{aligned} w_{t+1} &= w_t + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, w_t) = w_t + \alpha \cdot 1 \cdot (2\gamma - 1)w_t \cdot 1 \\ &= (1 + \alpha(2\gamma - 1))w_t \end{aligned}$$

Examples of Off-policy Divergence

Let us see more carefully what seen in the previous slide at the sequence of updates



The TD error on a transition between the two states is

$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t) = 0 + \gamma 2w_t - w_t = (2\gamma - 1)w_t$$

and the off-policy semi-gradient TD(0) update is

$$\begin{aligned} w_{t+1} &= w_t + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, w_t) = w_t + \alpha \cdot 1 \cdot (2\gamma - 1)w_t \cdot 1 \\ &= (1 + \alpha(2\gamma - 1))w_t \end{aligned}$$

Note that the importance sampling ratio, ρ_t , is 1 on this transition because there is only one action available from the first state, so its probabilities of being taken under the target and behavior policies must both be 1.

Examples of Off-policy Divergence

$$w_{t+1} = (1 + \alpha(2\gamma - 1))w_t$$

If $1 + \alpha(2\gamma - 1) > 1$ then the system is unstable and w will go to positive or negative infinite depending on its initial value

Examples of Off-policy Divergence

$$w_{t+1} = (1 + \alpha(2\gamma - 1))w_t$$

If $1 + \alpha(2\gamma - 1) > 1$ then the system is unstable and w will go to positive or negative infinite depending on its initial value



whenever $\gamma > 0.5$

Examples of Off-policy Divergence

$$w_{t+1} = (1 + \alpha(2\gamma - 1))w_t$$

If $1 + \alpha(2\gamma - 1) > 1$ then the system is unstable and w will go to positive or negative infinite depending on its initial value



whenever $\gamma > 0.5$

Note that stability does not depend on the specific step size, as long as $\alpha > 0$

Examples of Off-policy Divergence

$$w_{t+1} = (1 + \alpha(2\gamma - 1))w_t$$

If $1 + \alpha(2\gamma - 1) > 1$ then the system is unstable and w will go to positive or negative infinite depending on its initial value



whenever $\gamma > 0.5$

Note that stability does not depend on the specific step size, as long as $\alpha > 0$


Key to this example is that the one transition occurs repeatedly without w being updated on other transitions. *This is possible under off-policy training because the behavior policy might select actions on those other transitions which the target policy never would.* For these transitions, ρ_t would be zero and no update would be made.

Examples of Off-policy Divergence

This simple example communicates much of the reason why off-policy training can lead to divergence, but it is not completely convincing because it is not complete—it is just a fragment of a complete MDP.

Can there really be a complete system with instability? A simple complete example of divergence is *Baird's counterexample*.

Examples of Off-policy Divergence



Seven state, two-action MDP

The dashed action takes the system to one of the six upper states with equal probability, whereas the solid action takes the system to the seventh state


$$\pi(\text{solid}|\cdot) = 1$$

$$b(\text{dashed}|\cdot) = 6/7$$

$$b(\text{solid}|\cdot) = 1/7$$

$$\gamma = 0.99$$

Examples of Off-policy Divergence



The target policy π always takes the solid action and so the on-policy distribution is concentrated in the seventh state

Seven state, two-action MDP

The dashed action takes the system to one of the six upper states with equal probability, whereas the solid action takes the system to the seventh state


$$\pi(\text{solid}|\cdot) = 1$$

$$b(\text{dashed}|\cdot) = 6/7$$

$$b(\text{solid}|\cdot) = 1/7$$

$$\gamma = 0.99$$

Examples of Off-policy Divergence



The target policy π always takes the solid action and so the on-policy distribution is concentrated in the seventh state

Seven state, two-action MDP

The dashed action takes the system to one of the six upper states with equal probability, whereas the solid action takes the system to the seventh state

$$\pi(\text{solid}|\cdot) = 1$$


$$b(\text{dashed}|\cdot) = 6/7$$

$$b(\text{solid}|\cdot) = 1/7$$

$$\gamma = 0.99$$


The reward is zero on all transitions

Examples of Off-policy Divergence



Consider estimating the state-value under the linear parameterization indicated by the expression shown in each state circle.


Examples of Off-policy Divergence



Consider estimating the state-value under the linear parameterization indicated by the expression shown in each state circle.

For example, the estimated value of the leftmost state is $2w_1 + w_8$, where the subscript corresponds to the component of the overall weight vector $w \in \mathbb{R}^8$; this corresponds to a feature vector for the first state being $x(1) = (2, 0, 0, 0, 0, 0, 0, 1)^T$

Examples of Off-policy Divergence




Consider estimating the state-value under the linear parameterization indicated by the expression shown in each state circle.

For example, the estimated value of the leftmost state is $2w_1 + w_8$, where the subscript corresponds to the component of the overall weight vector $w \in \mathbb{R}^8$; this corresponds to a feature vector for the first state being $x(1) = (2, 0, 0, 0, 0, 0, 0, 1)^T$

The reward is zero on all transitions, so the true value function is $v_\pi(s) = 0$, for all s , which can be exactly approximated if $\mathbf{w} = 0$.

Examples of Off-policy Divergence


If we apply semi-gradient TD(0) $\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, \mathbf{w}_t)$ to this problem then the weights diverge to infinity



The instability occurs for any positive step size, no matter how small

Examples of Off-policy Divergence

If we apply semi-gradient TD(0) $\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \rho_t \delta_t \nabla \hat{v}(S_t, \mathbf{w}_t)$ to this problem then the weights diverge to infinity



The instability occurs for any positive step size, no matter how small

The example shows that even the *simplest combination of bootstrapping and function approximation can be unstable* if the updates are not done according to the on-policy distribution.

Examples of Off-policy Divergence

There are also counterexamples similar to Baird's showing divergence for Q-learning.

Examples of Off-policy Divergence

There are also counterexamples similar to Baird's showing divergence for Q-learning.

This is cause for concern because otherwise Q-learning has the best convergence guarantees of all control methods. Considerable effort has gone into trying to find a remedy to this problem or to obtain some weaker, but still workable, guarantee.

Examples of Off-policy Divergence

There are also counterexamples similar to Baird's showing divergence for Q-learning.

This is cause for concern because otherwise Q-learning has the best convergence guarantees of all control methods. Considerable effort has gone into trying to find a remedy to this problem or to obtain some weaker, but still workable, guarantee.

For example, it may be possible to guarantee convergence of Q-learning as long as the *behavior policy is sufficiently close to the target policy*, for example, when it is the ϵ -greedy policy. To the best of our knowledge, Q-learning has never been found to diverge in this case, but there has been no theoretical analysis.

The deadly triad

The danger of instability and divergence arises whenever we combine all of the following three elements

The deadly triad

The danger of instability and divergence arises whenever we combine all of the following three elements

Function approximation A powerful, scalable way of generalizing from a state space much larger than the memory and computational resources (e.g., linear function approximation or artificial neural networks).

The deadly triad

The danger of instability and divergence arises whenever we combine all of the following three elements

Function approximation A powerful, scalable way of generalizing from a state space much larger than the memory and computational resources (e.g., linear function approximation or artificial neural networks).

Bootstrapping Update targets that include existing estimates (as in TD methods) rather than relying exclusively on actual rewards and complete returns (as in MC methods).

The deadly triad

The danger of instability and divergence arises whenever we combine all of the following three elements

Function approximation A powerful, scalable way of generalizing from a state space much larger than the memory and computational resources (e.g., linear function approximation or artificial neural networks).

Bootstrapping Update targets that include existing estimates (as in TD methods) rather than relying exclusively on actual rewards and complete returns (as in MC methods).

Off-policy training Training on a distribution of transitions other than that produced by the target policy. Sweeping through the state space and updating all states uniformly does not respect the target policy and is an example of off-policy training.

The deadly triad

If any two elements of the deadly triad are present, but not all three, then instability can be avoided. It is natural, then, to go through the three and see if there is any one that can be given up.

The deadly triad

If any two elements of the deadly triad are present, but not all three, then instability can be avoided. It is natural, then, to go through the three and see if there is any one that can be given up.

Of the three, *function approximation most clearly cannot be given up.* We need methods that scale to large problems and to great expressive power

The deadly triad

Doing without bootstrapping is possible, at the cost of computational and data efficiency. Perhaps most important are the losses in computational efficiency. Monte Carlo (non-bootstrapping) methods require memory to save everything that happens between making each prediction and obtaining the final return, and all their computation is done once the final return is obtained.

With bootstrapping and eligibility traces data can be dealt with when and where it is generated, then need never be used again. The savings in communication and memory made possible by bootstrapping are great

The losses in data efficiency by giving up bootstrapping are also significant (faster learning)

The deadly triad

Off-policy learning; can we give that up? On-policy methods are often adequate. For model-free reinforcement learning, one can simply use Sarsa rather than Q-learning. *Using off-policy learning is considered as an appealing convenience but not a necessity.*

The deadly triad

Off-policy learning; can we give that up? On-policy methods are often adequate. For model-free reinforcement learning, one can simply use Sarsa rather than Q-learning. *Using off-policy learning is considered as an appealing convenience but not a necessity.*

However, off-policy learning is essential to other anticipated use cases, cases that we do not discuss in this class but may be important to the larger goal of creating a powerful intelligent agent.

In these use cases, the agent learns not just a single value function and single policy, but large numbers of them in parallel.

There are many target policies, and thus the one behavior policy cannot equal all of them

Gradient Temporal-Difference learning

- TD does not follow the gradient of *any* objective function
- This is why TD can diverge when off-policy or using non-linear function approximation
- **Gradient TD** follows true gradient of projected Bellman error

On/Off-Policy	Algorithm	Table Lookup	Linear	Non-Linear
On-Policy	MC	✓	✓	✓
	TD	✓	✓	✗
	Gradient TD	✓	✓	✓
Off-Policy	MC	✓	✓	✓
	TD	✓	✗	✗
	Gradient TD	✓	✓	✓

Convergence of control algorithms

Algorithm	Table Lookup	Linear	Non-Linear
Monte-Carlo Control	✓	(✓)	✗
Sarsa	✓	(✓)	✗
Q-learning	✓	✗	✗
Gradient Q-learning	✓	✓	✗

(✓) = chatters around near-optimal value function

Material

The material illustrated in this second part of the lecture can be found in the first three paragraphs of Chapter 11.

In the book there are more descriptions and thoughts. For the final examination pay attention to what seen in the slides used for the lecture (as usual concentrate on the algorithms!)